

Politecnico di Torino

Master's Degree in Computer Engineering

Migrating Legacy Services to a Cloud-Native Architecture: A Case Study on Outsight Cloud



outsight

Academic Supervisors: Prof. Fulvio Giovanni Risso, Stefano Galantino

Company Supervisors: Kasper Rutten, Anatolii Kostrygin

Candidate: Giovanni Russo

Matricola: s343424

Academic Year 2025–2026

Abstract

Legacy systems often evolve through years of incremental development, resulting in architectures that become difficult to scale, maintain, and modernize. This thesis investigates the migration of Oversight Cloud, a production platform for managing LiDAR-related workflows and assets, from a legacy service-oriented architecture toward a cloud-native environment on AWS.

Initially, the platform was deployed on virtual machines provisioned on the DigitalOcean cloud and orchestrated through Docker Compose, which managed the life-cycle of containerized application services directly on each VM without a dedicated container orchestration platform. Core business entities were stored in Airtable, which served as the primary operational database, while additional external services supported storage, documentation, and notification workflows. Over time, this architecture progressively introduced scalability, consistency, maintenance, and operational-governance challenges.

The thesis proposes and evaluates an incremental migration strategy based on Infrastructure as Code, managed cloud services, and progressive coexistence between legacy and target components. Particular attention is given to the transition from Airtable to PostgreSQL, the migration from VM-based deployments to managed container services, the adoption of Terraform-driven infrastructure provisioning, and the introduction of validation mechanisms designed to reduce migration risk while preserving service continuity. The proposed approach combines architectural decision-making with technical and operational evaluation criteria, including reproducibility, maintainability, deployment reliability, observability, and cost sustainability. To reduce migration risk and preserve system continuity, the project adopts incremental validation practices, automated consistency checks, and reproducible deployment workflows across isolated environments. The resulting architecture demonstrates

how cloud-native modernization can be achieved incrementally in a real industrial context through controlled evolution, continuous validation, and risk-aware design practices. Beyond the specific case study, the migration principles and engineering practices discussed throughout the thesis are applicable to broader legacy modernization scenarios, offering transferable insights for organizations facing similar cloud adoption challenges.

Acknowledgements

I would like to express my sincere gratitude to Oversight for providing me with the opportunity to complete this internship and work on challenging and meaningful projects. I am particularly thankful to my supervisor and the entire development team for their guidance, support, and patience throughout this journey.

I would also like to thank the Politecnico di Torino for the academic foundation that made this work possible, and my thesis advisor for valuable feedback and encouragement.

Finally, I am deeply grateful to my family and friends for their unwavering support during my studies.

Contents

Abstract	I
Acknowledgements	III
List of Figures	VIII
List of Tables	IX
1 Introduction	2
1.1 Industrial Context and Legacy Constraints	2
1.2 Problem Statement	3
1.3 Thesis Structure	4
2 Background and Related Work	6
2.1 Cloud-Native Principles for Legacy Modernization	6
2.1.1 Incremental Delivery and Operability	7
2.1.2 Infrastructure as Code	7
2.2 Migration Strategies and Patterns	8
2.2.1 Cloud Migration Fundamentals	8
2.2.2 System Coexistence in Incremental Migration	9
2.2.3 Decision-Driven Migration	10
2.3 IAM Modernization in Legacy Systems	10
2.3.1 From Application-Coupled Identity to Standard Protocols . .	11
2.3.2 Open-Source and Managed IAM Trade-offs	11
2.4 Limitations in Existing Approaches	12
2.4.1 Fragmented Treatment of Technical and Economic Criteria . .	12

2.4.2	Limited Evidence on Incremental Validation	12
2.5	Positioning of This Thesis	13
3	Case Study Baseline (As-Is System)	14
3.1	Current Architecture and Data Flow	14
3.2	Airtable as Primary Data Storage	15
3.3	Limitations of Airtable	17
3.3.1	Data Modeling and Consistency	17
3.3.2	Limited Relational Capabilities	18
3.3.3	Operational Inefficiency	18
3.3.4	Scalability, Coupling, and Cost	18
3.4	Motivation for Migration to PostgreSQL	19
4	Migration Methodology and Decision Framework	20
4.1	Migration Goals and Success Criteria	20
4.2	Candidate Solutions Considered	21
4.2.1	Infrastructure and Runtime Options	21
4.2.2	Container Orchestration: Kubernetes vs ECS Fargate	22
4.2.3	Alternative Cloud Platforms Considered	23
4.2.4	Data-Layer Migration Options	25
4.2.5	Identity Evolution Options	25
4.3	Technical and Economic Evaluation Criteria	26
4.3.1	Technical Criteria	26
4.3.2	Economic and Operational Criteria	27
4.4	Selected Strategy and Rationale	27
4.5	Risk Management and Coexistence Plan	29
5	Target Architecture and Solution Design (To-Be)	30
5.1	Architecture Overview	30
5.2	Client and Delivery Layer	32
5.2.1	Browser	32
5.2.2	Amazon CloudFront	32
5.3	Load Balancing Layer	32
5.3.1	Application Load Balancer	32

5.4	Application Layer	32
5.4.1	ECS Fargate Runtime	32
5.4.2	Frontend Service (Vue)	33
5.4.3	Backend Service (Node.js API)	33
5.5	Managed Services Layer	33
5.5.1	PostgreSQL	33
5.5.2	Redis	33
5.5.3	Amazon S3	34
5.5.4	Amazon Cognito	34
5.5.5	Amazon SES	34
5.5.6	Amazon SNS	34
5.6	Network and Security Topology	34
5.7	Architectural Benefits	35
6	Migration Implementation	38
6.1	Incremental Execution Plan	38
6.2	Infrastructure as Code and Environment Separation	40
6.2.1	Infrastructure as Code (IaC)	40
6.2.2	Terraform-Driven Provisioning	41
6.2.3	Terraform Workflow: Write, Plan, Apply	41
6.2.4	State Management and Team Collaboration	42
6.2.5	Rationale for Adopting Terraform	42
6.2.6	Environment Governance	43
6.3	Infrastructure Migration Execution	43
6.3.1	Target Environment Definition and Terraform Provisioning	43
6.3.2	Routing, Application Adaptation, and Container Deployment	44
6.3.3	Storage, TLS, and Content Delivery	44
6.3.4	Environment Validation and Controlled Service Transition	45
6.4	Data Migration Tooling and Validation	46
6.4.1	Relational Foundation	46
6.4.2	Replication and Validation Scripts	46
7	Evaluation and Discussion	48
7.1	Benchmark Methodology and Demand Profile	48

7.1.1	Measurement Scope	48
7.1.2	Workload and Demand Definition	50
7.1.3	As-Is vs To-Be Benchmark Setup	51
7.2	Cost and TCO Analysis	51
7.2.1	Current Architecture Costs (As-Is)	52
7.2.2	Target Architecture Costs (To-Be)	52
7.2.3	Total Cost of Ownership (TCO)	53
7.3	Latency and Performance Analysis	55
7.3.1	As-Is Latency Measurements	55
7.3.2	To-Be Latency Measurements	56
7.3.3	Comparative Results	57
7.4	Operational Cost and Efficiency	59
7.4.1	As-Is Operational Effort	59
7.4.2	To-Be Operational Effort	60
7.4.3	Operational Cost Comparison	62
7.4.4	External Validation of the Operational Estimates	63
8	Conclusions and Future Work	64
8.1	Key Outcomes	64
8.2	Limitations and Open Questions	65
8.3	Future Work	66
	Bibliography	68

List of Figures

3.1	Outsight Cloud baseline architecture before the migration.	16
5.1	Final AWS architecture adopted for Outsight Cloud.	31
6.1	Incremental execution flow from the legacy baseline to the target AWS cloud-native architecture	39
7.1	Estimated steady-state AWS cost by service (EUR/month).	53
7.2	Monthly cost by category: Legacy vs. AWS (EUR/month). Cate- gories absent from the legacy platform are shown as zero.	54
7.3	Mean latency comparison, Legacy vs. AWS, per HTTP timing com- ponent, with 95% CI error bars.	57
7.4	Box plots of latency distributions per HTTP timing component (100 samples per environment).	58
7.5	ECDFs of latency per HTTP timing component, with p90/p95/p99 reference lines.	58

List of Tables

4.1	Comparative assessment of cloud provider alternatives for Oversight Cloud migration	25
7.1	HTTP timing metrics extracted from <code>curl</code> and their meaning. All values are recorded in seconds by <code>curl</code> and converted to milliseconds during post-processing.	49
7.2	Benchmark configuration parameters used for all latency measurements.	50
7.3	Infrastructure comparison between the As-Is (Legacy) and To-Be (AWS) benchmark targets. Both environments expose the same application over HTTPS.	51
7.4	Monthly cost comparison between the Legacy and AWS platforms. . .	53
7.5	Descriptive statistics for the Home endpoint benchmark ($N = 100$ requests per environment). All timing values in milliseconds (ms). . .	55
7.6	Operational effort comparison between the Legacy and AWS architectures. Estimates are in engineering hours per month.	62

Chapter 1

Introduction

1.1 Industrial Context and Legacy Constraints

In recent years, cloud computing has significantly transformed the way software systems are designed, deployed, and maintained. Modern cloud platforms provide scalable infrastructure, managed services, automated deployment capabilities, and operational flexibility that are difficult to achieve with traditional self-managed environments. As a consequence, many organizations are progressively modernizing legacy applications to adopt cloud-native architectures.

However, migrating an existing production platform to the cloud is often far from a straightforward process. Legacy systems are typically the result of years of incremental development shaped by evolving business requirements, operational constraints, and technological decisions. Over time, these systems often accumulate architectural dependencies and operational practices that complicate scalability, maintainability, and long-term evolution.

This thesis analyzes the migration of Outsight Cloud (OC), a platform used to manage LiDAR-related workflows, simulations, organizations, sites, and associated resources. Before the migration effort, the platform relied on a virtual-machine-based infrastructure hosted on DigitalOcean Droplets and orchestrated through Docker Compose, integrating several external services, most notably Airtable for core data management. This architecture enabled rapid development and operational flexibility during the early stages of the product lifecycle, allowing new features to be delivered with limited infrastructure overhead. As platform usage and

integrations expanded, however, this approach began to show signs of architectural strain.

1.2 Problem Statement

The engineering problem addressed in this work concerns the migration of Oversight Cloud from a legacy architecture toward a cloud-native environment under real industrial and operational constraints.

Over time, the platform became strongly dependent on external services and platform-specific integrations, especially *Airtable*, which was deeply embedded in application workflows and operational processes. The absence of a fully relational data model limited consistency guarantees and complicated the management of relationships between entities. Operational procedures increasingly relied on manual or partially automated workflows, generating inefficiencies and increasing the risk of human error. In parallel, the original deployment model became progressively harder to scale and maintain as infrastructure requirements and operational complexity continued to grow.

To address these limitations, the company initiated a cloudification process aimed at redesigning the platform architecture and migrating toward a managed cloud environment. The selection of the target platform and architectural strategy was the result of a structured evaluation of candidate solutions against technical, operational, and economic criteria. The main engineering objective was to regain explicit architectural ownership over core data and infrastructure assets through Infrastructure as Code practices, the migration from *Airtable* to *PostgreSQL*, the introduction of managed cloud services, and the modernization of deployment and operational workflows. A central challenge was defining an execution sequence capable of reducing technical and operational risk before final cutover, relying on validation workflows, reproducible infrastructure management, and controlled transition mechanisms.

Beyond addressing these architectural limitations, the migration initiative was driven by a set of concrete business and engineering objectives:

- Reduce operational costs by replacing *Airtable* and self-managed infrastructure with managed cloud services.

- Improve system performance and data access efficiency through the adoption of PostgreSQL as the primary relational datastore.
- Reduce operational overhead by automating infrastructure provisioning, deployment, and maintenance through Infrastructure as Code practices.
- Minimize migration risk and service disruption through incremental coexistence, validation workflows, and controlled transition mechanisms.

1.3 Thesis Structure

This thesis is organized as follows:

Chapter 2 introduces the conceptual and related-work foundations on cloud-native modernization, migration patterns, and IAM evolution in legacy systems.

Chapter 3 describes the case-study baseline (as-is architecture), including technical debt and cost-related migration drivers.

Chapter 4 presents the migration methodology and decision framework, including candidate options and selection rationale.

Chapter 5 defines the target architecture (to-be), its main components, and design trade-offs.

Chapter 6 details the incremental migration implementation foundations, with focus on IaC, data migration tooling, and validation pipeline.

Chapter 7 presents the evaluation of the migration in terms of infrastructure costs and total cost of ownership, performance, and operational effort.

Chapter 8 concludes the thesis by summarizing contributions and outlining future work.

Chapter 2

Background and Related Work

This chapter provides the conceptual background required to frame the migration problem and positions the thesis with respect to related work.

2.1 Cloud-Native Principles for Legacy Modernization

Cloud-native modernization represents a broader approach to evolving software systems, combining architectural changes with new practices for managing and operating applications. In the context of legacy migration, the goal is to establish a more structured and sustainable foundation that can support the progressive evolution of the system. This involves improving how infrastructure, application components, and operational processes are defined, deployed, and managed throughout the migration. These principles become particularly important when architectural changes affect production workflows, persistent data, and operational procedures.

In *Building Microservices*, Newman frames modern service-oriented architectures around independently deployable services with explicit boundaries and hidden internal state [1]. This principle is relevant to legacy modernization because it separates external contracts from internal implementation details, reducing the risk that changes in one component propagate unexpectedly across the platform.

Modern cloud-native systems are typically designed around automation-oriented practices and managed services that simplify infrastructure operations and improve

scalability. However, migrating a legacy platform toward such architectures requires careful consideration of production constraints, compatibility requirements, and long-term maintainability.

2.1.1 Incremental Delivery and Operability

A recurring principle in modernization literature is the preference for incremental delivery over full-system replacement. Incremental migration reduces operational risk, enables earlier feedback cycles, and allows validation activities to be performed progressively throughout the transition process.

Newman argues against big-bang rewrites, emphasizing that large simultaneous replacements tend to concentrate risk and delay feedback [1]. Decomposing migration work into smaller steps makes it possible to start from lower-risk components, build confidence progressively, and validate architectural changes before irreversible decisions are made.

To be effective, incremental delivery requires infrastructure automation, environment consistency, reproducible deployment workflows, and testable interfaces. In industrial systems, these practices make architectural changes observable, reversible, and easier to validate before they affect critical workflows.

2.1.2 Infrastructure as Code

Infrastructure as Code (IaC) is a key enabler for cloud-native modernization because it transforms infrastructure configuration into versioned and reproducible artifacts. Instead of relying on manual provisioning procedures, infrastructure resources are defined declaratively and managed through automated workflows.

Newman identifies Infrastructure as Code as a core deployment principle because infrastructure configuration stored in source control allows environments to be recreated deterministically and shared across teams [1]. In *Infrastructure as Code: Dynamic Systems for the Cloud Age*, Morris further identifies three core practices for effective IaC adoption: defining everything as code, continuously testing and delivering all work in progress, and building small loosely coupled pieces that can be changed independently [2]. Together, these practices improve traceability, reproducibility, reviewability, and consistency across development and production

environments. In migration projects, IaC also simplifies environment replication, deployment governance, and rollback management, reducing operational risks during infrastructure evolution.

2.2 Migration Strategies and Patterns

Cloud migration projects can follow different strategies depending on system complexity, operational requirements, and modernization goals. Common approaches include rehosting, replatforming, refactoring, and selective rebuilding. In practice, enterprise migration projects often combine multiple strategies across subsystems according to their coupling level, criticality, migration effort, and risk exposure.

2.2.1 Cloud Migration Fundamentals

Cloud migration is the process of moving applications, data, and infrastructure resources from traditional environments to cloud-based platforms [3]. In practice, this transition is executed progressively through staged rollout strategies and validation activities designed to minimize disruption.

Historically, many organizations operated software systems on self-managed infrastructures hosted in local data centers or virtualized environments. While these architectures provided direct control over resources, they also introduced significant operational complexity related to infrastructure maintenance, capacity planning, backup management, monitoring, and service availability.

For this reason, cloud migration planning must evaluate not only where workloads will run, but also how infrastructure responsibilities, automation practices, reliability requirements, and operational governance will change in the target environment.

Reference guidance highlights several recurring migration drivers and benefits [3], including cost optimization, elastic scalability, managed operational services, improved infrastructure reliability, and access to modern distributed delivery models.

Cloud migration strategies are generally classified into three major categories [3]:

- **Rehosting**, which moves workloads to the cloud with minimal architectural

changes.

- **Replatforming**, which introduces selective adaptations to better exploit cloud capabilities while preserving the overall application architecture.
- **Refactoring**, which redesigns parts of the system to align with cloud-native principles and improve long-term scalability and maintainability.

Migration programs are often organized into assessment, preparation, and migration phases [3]. These phases support technical discovery, environment preparation, pilot execution, and incremental rollout activities.

In practice, real-world migration projects often combine multiple strategies across different system components. Newman also stresses that decomposition and modernization choices should be driven by the desired outcome rather than by technical convenience alone [1]. The Outsight Cloud migration can be primarily classified as a replatforming effort complemented by selective refactoring activities.

The migration preserved the core application architecture and business functionality while replacing key infrastructure and platform dependencies. Examples of replatforming include the transition from DigitalOcean-hosted workloads to AWS managed services and the adoption of ECS Fargate as the runtime environment.

At the same time, selected components required refactoring to support the migration objectives. These changes included the transition from Airtable to PostgreSQL, the introduction of Infrastructure as Code practices through Terraform, and adaptations to deployment and routing workflows required by the target cloud-native architecture.

2.2.2 System Coexistence in Incremental Migration

For production systems that cannot tolerate service interruption, coexistence between legacy and target components becomes a critical architectural requirement. During migration phases, both systems frequently need to operate simultaneously while compatibility and behavioral consistency are progressively validated.

Newman describes the *strangler fig pattern* as a technique for wrapping legacy systems incrementally and redirecting calls to new implementations without replacing the entire system at once [1]. He also discusses parallel-run approaches, where

old and new implementations execute side by side and their results are compared before final cutover [1].

Morris further describes the *expand and contract* pattern as a structured technique for changing live infrastructure without disrupting running services. The pattern involves first adding new resources alongside existing ones, migrating consumers incrementally, and only then removing the legacy components [2]. This approach avoids the risks of big-bang replacements and keeps the system in a consistent, testable state throughout the transition.

This coexistence model introduces additional engineering complexity, including contract compatibility, dual-backend validation, synchronization workflows, rollback capabilities, and progressive switchover mechanisms.

2.2.3 Decision-Driven Migration

Related work frequently describes migration patterns from a purely technological perspective, focusing primarily on target architectures and adopted cloud services. In industrial contexts, however, migration decisions must also consider operational constraints, organizational impact, recurring costs, maintainability, and migration risk.

For this reason, migration strategies should be supported by explicit decision frameworks capable of combining technical and economic rationale. This aspect is particularly relevant in large-scale modernization projects where architectural choices affect both infrastructure sustainability and long-term operational governance.

2.3 IAM Modernization in Legacy Systems

Identity and Access Management (IAM) modernization is a recurrent challenge in legacy platforms. Historically, many systems implemented authentication and authorization logic directly inside application workflows and persistence layers, generating tightly coupled identity-management models that become difficult to evolve and maintain over time.

In AWS-based environments, *AWS Identity and Access Management* (IAM) is the web service that governs access to cloud resources: it defines who is authenticated

and who is authorized to use them, and associates *principals* (such as IAM users, roles, federated identities, or applications) with permission policies that determine which operations are allowed on account resources [4]. Access is granted only after an authorization request succeeds against the policies in effect for that principal and target service [4].

Modern cloud-native architectures instead rely on standardized protocols and dedicated identity services to improve interoperability, scalability, and security.

2.3.1 From Application-Coupled Identity to Standard Protocols

Legacy systems frequently embed authentication assumptions directly into business logic and data models. Migrating toward standards-based IAM therefore requires decoupling identity management from application-specific implementations and re-defining authorization boundaries.

Protocols such as OAuth 2.0 and OpenID Connect (OIDC) provide stronger interoperability and security guarantees than ad-hoc authentication approaches. However, introducing standardized IAM mechanisms into existing systems also requires compatibility validation with existing workflows, clients, and operational procedures.

2.3.2 Open-Source and Managed IAM Trade-offs

The choice between open-source IAM solutions and managed identity platforms affects operational complexity, integration flexibility, infrastructure ownership, and long-term maintenance costs.

Managed IAM services simplify operational governance by reducing infrastructure-management responsibilities and providing built-in support for authentication workflows, user lifecycle management, and security integration. Open-source solutions, on the other hand, generally provide greater customization flexibility and infrastructure control at the cost of increased operational responsibility.

In migration projects, these trade-offs also influence how progressively existing services can adopt modern authentication flows while preserving backward compatibility.

2.4 Limitations in Existing Approaches

The literature provides strong conceptual foundations on cloud migration patterns and IAM technologies. However, industrial guidance often remains fragmented across separate technical dimensions.

2.4.1 Fragmented Treatment of Technical and Economic Criteria

Technical architecture decisions are frequently discussed without detailed consideration of operational costs, infrastructure governance, or long-term maintenance overhead, even though sustainable migration depends on aligning technical choices with those same constraints.

As a consequence, migration projects require evaluation frameworks capable of combining scalability, maintainability, operational sustainability, and economic impact within a unified decision process.

2.4.2 Limited Evidence on Incremental Validation

Many migration case studies focus primarily on final target architectures while providing limited evidence regarding coexistence management and intermediate validation mechanisms.

In industrial environments, migration reliability often depends on validation practices such as dual-backend testing, data-consistency verification, environment-separated deployment workflows, and incremental rollout strategies. Morris emphasizes that continuously testing infrastructure changes as they are developed rather than waiting until the system is complete is essential for catching errors early and reducing the cost of correction [2]. These aspects are essential for reducing migration risk during long transition phases.

2.5 Positioning of This Thesis

This thesis addresses the above limitations through a decision-oriented migration case study focused on incremental architectural evolution under real industrial constraints.

The contribution of the work is not limited to the description of a target AWS architecture. Instead, it analyzes the migration of a production platform from a legacy DigitalOcean-based environment toward a cloud-native architecture built on managed AWS services. The thesis also emphasizes the reasoning process that connects baseline system limitations, migration decisions, infrastructure automation, data-layer transition planning, and validation workflows.

Particular attention is given to Infrastructure as Code adoption, migration reproducibility, environment separation, and validation-oriented migration practices designed to reduce technical and operational risk.

Although the work is centered on Outsight Cloud, the migration principles and engineering practices discussed throughout the thesis are applicable to broader legacy modernization scenarios involving progressive cloud-native adoption.

Chapter 3

Case Study Baseline (As-Is System)

This chapter describes the initial state of the Oversight Cloud platform before the migration strategy was formalized. The objective is to establish the architectural and operational baseline against which migration decisions and modernization efforts are evaluated.

3.1 Current Architecture and Data Flow

Before the migration to AWS, Oversight Cloud relied on a containerized architecture designed to support rapid feature delivery and operational flexibility during the early stages of product development.

The platform managed LiDAR-related workflows and assets through a web application and backend API stack composed of several interconnected services. The main components of the legacy architecture were:

- **Vue 3 SPA Frontend** responsible for the user interface and interaction layer.
- **Node.js/Express Backend API** implementing business logic, workflow orchestration, and integrations with external services.
- **Redis** used for session management and temporary application state.

- **Docker Compose** used to orchestrate application containers on DigitalOcean Droplets.
- **Traefik Reverse Proxy** providing HTTPS termination and request routing.
- **Airtable** acting as the primary operational datastore for users, organizations, sites, simulations, permissions, and role assignments.
- **S3-Compatible Object Storage** used for simulation files and asset storage.
- **GitBook** used as the platform documentation system.
- **Slack and Customer.io Integrations** supporting notifications and operational workflows.
- **DigitalOcean Droplets** providing the virtual-machine-based infrastructure hosting the application stack.

From an infrastructure perspective, the architecture followed a virtual-machine-centric model. Application services were deployed as long-running containers, with scaling, upgrades, deployment management, and operational maintenance handled directly by the engineering team. While this approach was sufficient during the initial stages of the platform, it increasingly exposed limitations related to scalability, operational governance, fault tolerance, and infrastructure reproducibility.

Although the stack enabled fast iteration and reduced initial infrastructure complexity during early product development, it combined self-managed virtual-machine infrastructure with multiple external service dependencies.

3.2 Airtable as Primary Data Storage

As shown in the component overview above, Airtable was more than a persistence layer: it acted simultaneously as operational datastore, workflow-management layer, and integration hub for several internal processes.

Airtable is a cloud-based platform that combines spreadsheet-like interfaces with lightweight database capabilities. Data is organized into bases, tables, records, and fields, allowing users to manage structured information through a graphical and




Figure 3.1. Oversight Cloud baseline architecture before the migration.

collaborative environment. Airtable also provides features such as views, forms, automations, integrations, and API access, making it particularly suitable for rapid prototyping and operational workflow management [5].

In Oversight Cloud, these capabilities were exercised through API access, operational forms, and integrations with external systems such as Slack-based notification workflows. During the early phases of the project, the accessibility of Airtable for non-technical stakeholders accelerated development and simplified collaboration between engineering and operations teams.

3.3 Limitations of Airtable

As Oversight Cloud became more complex, the limitations of using Airtable as primary backend infrastructure became increasingly evident. The issues affected not only data consistency and backend maintainability, but also operational workflows, scalability, and long-term architectural evolution.

3.3.1 Data Modeling and Consistency

Relationships between entities such as organizations, sites, simulations, and permissions were represented through linked records and lookup fields rather than through explicit relational constraints.

As a consequence, consistency guarantees depended heavily on application-side logic instead of database-level enforcement mechanisms. Backend services frequently needed to normalize and reinterpret Airtable responses in order to maintain coherent internal representations.

In several cases, lookup fields exposed heterogeneous data structures depending on the access context. The same logical information could appear either as scalar values or arrays, increasing backend complexity and complicating synchronization and mapping logic.

This absence of strict schema enforcement progressively increased technical debt and reduced maintainability as the number of entities and relationships expanded.

3.3.2 Limited Relational Capabilities

Although Airtable supports basic linked-record relationships, it does not provide the full relational capabilities of SQL database systems such as PostgreSQL.

Advanced joins, indexing strategies, transactional consistency, complex filtering operations, and optimized querying mechanisms were significantly more limited. As business logic became more sophisticated, these limitations increasingly affected backend implementation complexity and query efficiency.

The platform therefore required growing amounts of application-side logic to compensate for missing relational features that would normally be handled directly by a relational DBMS.

3.3.3 Operational Inefficiency

Several operational workflows relied on manual interactions through Airtable forms and operational interfaces.

User onboarding and account-management procedures, for example, often required manual intervention from internal teams. This approach slowed operational processes, increased dependency on human intervention, and raised the probability of operational errors.

As the number of users, organizations, and managed simulations increased, these manual procedures became progressively harder to scale efficiently.

3.3.4 Scalability, Coupling, and Cost

Over time, Airtable evolved from a simple operational datastore into a critical architectural dependency tightly coupled with application workflows and operational procedures.

Business logic progressively became dependent on Airtable-specific concepts and APIs, reducing architectural flexibility and complicating future system evolution. This dependency concentration increased technical debt and limited the maintainability of the platform.

The economic impact also became significant. Recurring operational costs associated with Airtable usage created additional pressure to adopt a more scalable

and sustainable backend architecture, without depending on a third-party platform whose cost structure was difficult to control as usage increased.

3.4 Motivation for Migration to PostgreSQL

The migration from Airtable to PostgreSQL was initiated to address the architectural, operational, and economic limitations identified in the previous sections.

The objective was not simply to replace one storage technology with another, but to progressively transition from a semi-structured operational backend toward a maintainable and scalable relational architecture capable of supporting long-term platform evolution.

PostgreSQL introduced several advantages compared to the previous data management model. Explicit relational schemas based on primary and foreign keys enabled stronger consistency guarantees and improved integrity enforcement. Advanced querying capabilities, transactional support, and indexing mechanisms also simplified backend implementation and improved data management efficiency.

From an architectural perspective, the migration reduced dependence on Airtable-specific structures and allowed the platform to regain ownership of its data model and persistence logic.

Chapter 4

Migration Methodology and Decision Framework

This chapter presents the migration methodology adopted for the modernization of Outsight Cloud and explains the decision process that guided the selection of the target architecture and migration strategy. The focus is not limited to the technologies adopted, but also includes the technical, operational, and economic considerations that influenced the migration process.

4.1 Migration Goals and Success Criteria

Before defining the target architecture, a set of migration goals was identified in order to establish a structured decision framework and reduce solution bias during implementation activities.

The first objective was to avoid disruptive replacement strategies or prolonged service interruptions in production environments. This requirement translated into concrete success criteria: staged rollout capability, rollback readiness, validation of critical workflows before cutover, and separation between development and production migration activities.

Another major objective was the reduction of architectural coupling to legacy dependencies, particularly those related to Airtable-based operational workflows and platform-specific integrations. The migration also aimed to improve scalability, maintainability, and long-term architectural evolvability by adopting managed cloud

services and reproducible infrastructure-management practices.

Particular attention was also given to deployment reproducibility and environment governance. The migration strategy needed to support controlled infrastructure evolution across development and production environments while reducing manual operational activities.

In addition to technical goals, the migration process required explicit evaluation of operational sustainability and recurring infrastructure costs. Architectural decisions therefore had to balance technical quality, implementation complexity, operational overhead, and long-term maintainability.

From these objectives, several success criteria were derived. The migration process needed to support validation of critical application flows during the data-layer transition outlined in Chapter 3, provide traceable infrastructure definitions across environments, reduce dependency concentration around legacy services, and enable migration verification through tests, scripts, and CI-supported validation workflows.

4.2 Candidate Solutions Considered

The migration design process considered multiple alternatives for each architectural dimension, including infrastructure management, cloud-platform selection, data-layer evolution, and identity-management modernization.

Rather than adopting a purely technology-driven approach, the evaluation focused on identifying solutions capable of balancing migration risk, production stability, maintainability, and implementation effort.

4.2.1 Infrastructure and Runtime Options

Several infrastructure strategies were considered during the planning phase of the migration.

The first option consisted of preserving the existing deployment model while progressively hardening operational workflows and infrastructure management practices. This approach minimized short-term migration complexity but did not adequately address scalability limitations, operational governance issues, or long-term maintainability concerns.

A second option involved migrating toward managed cloud services combined with Infrastructure as Code practices. This strategy introduced a more structured cloud-native operational model while preserving incremental migration capabilities and reducing infrastructure-management overhead.

A third possibility was the adoption of a fully orchestration-oriented redesign from the beginning, for example through Kubernetes-centered infrastructure. Although this approach provided strong scalability and orchestration capabilities, it significantly increased migration complexity and operational overhead during the early transition phases.

4.2.2 Container Orchestration: Kubernetes vs ECS Fargate

During the infrastructure design phase, Kubernetes-based solutions were evaluated as a potential runtime platform for containerized workloads. Kubernetes is a portable, extensible open source platform for managing containerized workloads and services through declarative configuration and automation [6]. In production environments, it provides a framework to run distributed systems resiliently by handling scaling, failover, service discovery, load balancing, and controlled rollouts [6].

However, adopting Kubernetes would have introduced additional operational complexity that was not justified by the requirements of Oversight Cloud at the time of the migration. Managing Kubernetes clusters requires expertise in cluster administration, networking, upgrades, monitoring, capacity planning, and security governance. These responsibilities would have increased both implementation effort and long-term operational overhead.

This evaluation is consistent with Newman’s caution against premature platform complexity: small service landscapes do not necessarily require Kubernetes-level orchestration, especially when the operational cost of cluster management would distract from migration goals [1]. For smaller-scale migrations, simpler managed container platforms can provide a better balance between deployment flexibility and operational effort.

In contrast, Amazon ECS Fargate provides a managed container execution environment that eliminates the need to provision and maintain worker nodes. The service offers native integration with other AWS components, including Application Load Balancer, CloudWatch, IAM, and Auto Scaling, while preserving support for

Docker-based workloads already used by the existing platform.

From a migration perspective, ECS Fargate enabled a smoother transition from the legacy Docker Compose deployment model because applications could continue to be packaged as standard containers without requiring a complete redesign of the runtime environment.

Given the project objectives (incremental migration, reduced operational burden, faster adoption, and maintainability), ECS Fargate provided the most appropriate balance between scalability and operational simplicity. While Kubernetes remains a viable option for future platform evolution, its additional complexity was not justified by the scale and requirements of the current migration effort.

Given the project constraints and the need for controlled incremental evolution, the migration strategy favored managed cloud services combined with reproducible infrastructure provisioning practices, with ECS Fargate selected as the container runtime platform.

4.2.3 Alternative Cloud Platforms Considered

During the migration design phase, multiple cloud platform alternatives were evaluated before selecting the final target environment. The comparison focused on scalability, availability of managed services, Infrastructure as Code support, ecosystem maturity, operational complexity, and compatibility with the existing technical baseline.

The main alternatives considered were Amazon Web Services (AWS), Microsoft Azure, Google Cloud Platform (GCP), continuation on the existing DigitalOcean infrastructure, and self-managed on-premises infrastructure.

Amazon Web Services (AWS). AWS was considered the strongest candidate due to the maturity of its managed-service ecosystem and its compatibility with the migration objectives of Outsight Cloud.

Services such as ECS Fargate, RDS PostgreSQL, S3, Cognito, SES, and CloudFront provided native support for the architectural model targeted by the migration strategy. In addition, AWS offered strong Terraform integration, mature DevOps tooling, and operational scalability aligned with the company's broader cloudification strategy.

The existing operational knowledge inside the company was also already partially

aligned with AWS-based workflows, reducing migration overhead and simplifying infrastructure adoption.

Microsoft Azure. Microsoft Azure was also evaluated as a potential target environment. Azure provides a broad enterprise ecosystem with managed services covering compute, storage, networking, databases, and DevOps workflows.

Equivalent services for most components used in this project were available, including Azure Virtual Machines, Azure Kubernetes Service (AKS), Azure SQL Database, and Azure Blob Storage.

Google Cloud Platform (GCP). Google Cloud Platform was considered primarily because of its strengths in containerized workloads and Kubernetes-native operations.

Relevant services included Google Kubernetes Engine (GKE), Cloud SQL, Compute Engine, and Cloud Storage. GCP also provides strong capabilities in analytics and large-scale distributed infrastructure management.

Continuation on DigitalOcean. Another alternative consisted of preserving the existing infrastructure provider and progressively improving the deployment model without changing cloud platform. This approach would have minimized migration effort in the short term and leveraged existing operational knowledge. However, it would not have provided the same level of managed-service integration, infrastructure automation, and cloud-native capabilities offered by AWS. As a consequence, it was considered less suitable for supporting the long-term modernization objectives of the platform.

On-Premises Infrastructure. Another possible strategy consisted of continuing with self-managed infrastructure hosted on virtual machines or on-premises-style environments.

Although this approach offered direct control over infrastructure resources and configurations, it significantly increased operational responsibilities related to maintenance, monitoring, backup management, scalability planning, high availability, and infrastructure security.

This option conflicted with one of the primary migration goals: reducing operational burden through managed cloud services and automation-oriented infrastructure governance.

Table 4.1. Comparative assessment of cloud provider alternatives for Out-sight Cloud migration

Criterion	AWS	Azure	GCP	On-prem
Managed services fit	High	Medium	Medium	Low
IaC and automation	High	Medium	Medium	Low
Operational complexity	Low to medium	Medium	Medium	High
Ecosystem and maturity	Very high	High	High	Variable
Migration continuity with current baseline	High	Medium-low	Medium-low	Low

4.2.4 Data-Layer Migration Options

The limitations of the legacy datastore described in Chapter 3 shifted the evaluation away from the question of whether PostgreSQL should be adopted and toward the more practical question of *how* to transition toward the target relational model.

A first possibility was an immediate cutover strategy. Although this approach simplified long-term architecture, it introduced high migration risk and limited rollback flexibility.

A second option involved maintaining long-term dual-backend operation without defining a clear migration endpoint. While operationally safer in the short term, this approach risked increasing architectural complexity and prolonging dependency concentration around legacy systems.

The selected approach instead relied on progressive validation mechanisms and controlled switchover procedures. This model allowed data-transfer correctness, API behavior, and rollback readiness to be verified before the final transition to PostgreSQL.

4.2.5 Identity Evolution Options

The migration process also required evaluating alternative approaches for identity and authentication modernization.

One option consisted of preserving existing authentication contracts and postponing IAM modernization until after infrastructure migration. While simpler initially, this strategy risked extending architectural coupling to legacy authentication assumptions.

A second possibility involved introducing standards-based IAM mechanisms progressively while preserving compatibility with existing services and workflows. This approach enabled gradual adoption of modern authentication practices while reducing migration disruption.

A third option consisted of a complete redesign of authentication and authorization mechanisms from the beginning of the migration project. Although architecturally cleaner, this approach introduced excessive migration complexity and operational risk during early transition stages.

For these reasons, the migration strategy favored progressive IAM modernization based on compatibility-preserving authentication changes and staged adoption of standards-based flows.

4.3 Technical and Economic Evaluation Criteria

Each candidate solution was evaluated against a common set of technical and operational criteria in order to support structured architectural decision-making.

4.3.1 Technical Criteria

Technical evaluation focused on aspects directly related to migration feasibility, production stability, and long-term maintainability.

Particular importance was given to backward compatibility with existing workflows and service behaviors, because API contracts, authentication assumptions, and data-access paths had to remain predictable during staged changes. Security considerations were also central, especially regarding readiness for standards-based IAM integration and cloud-native infrastructure governance.

Deployment reproducibility represented another key criterion. Candidate solutions needed to support versioned infrastructure definitions, automated provisioning workflows, and consistent environment management across development and production stages.

Migration testability was also considered essential. The selected architecture needed to support incremental validation workflows, rollback readiness, and automated verification practices capable of reducing migration risk.

Finally, operational observability and troubleshooting support were evaluated in order to ensure visibility during rollout phases.

4.3.2 Economic and Operational Criteria

In addition to technical aspects, the evaluation process also considered operational sustainability and recurring infrastructure costs.

Particular attention was given to reducing operational overhead associated with infrastructure maintenance and manual workflows. Candidate solutions were also evaluated according to their ability to reduce dependency concentration around high-cost external services and simplify long-term operational governance.

Maintainability, scalability, and automation capabilities were therefore analyzed together with recurring infrastructure costs and migration effort distribution.

4.4 Selected Strategy and Rationale

The selected strategy was an incremental migration centered on managed cloud services, Infrastructure as Code practices, and independently verifiable migration steps.

The migration involved replacing the legacy infrastructure with a cloud-native architecture built on AWS managed services. The selected strategy prioritized progressive modernization of infrastructure, deployment workflows, and data-management components through changes that could be validated separately.

AWS-managed services were adopted for runtime execution, routing, storage, authentication, and observability, with ECS Fargate selected over Kubernetes as the container runtime platform (see Section 4.2.2). Terraform became the primary infrastructure control layer responsible for reproducible provisioning and environment governance.

For the data layer, the strategy applied the controlled transition model defined in Section 4.2.4, building on the PostgreSQL target identified in Chapter 3.

AWS was selected because it provided the most complete alignment with the technical and operational requirements of the project. In particular, the platform offered mature Terraform integration, strong scalability characteristics, managed services aligned with the target architecture (including RDS PostgreSQL), and compatibility with the broader cloudification strategy already adopted by the company.

AWS was also considered highly developer-friendly, with broad customizability enabled by Lambda-based logic and service integrations that fit the operating model of a SaaS company requiring implementation flexibility.

The main trade-off acknowledged during the evaluation is that AWS can become expensive under specific scaling and usage conditions. However, for this project, the balance between service quality, customization freedom, and migration continuity remained favorable to AWS.

The selected approach balanced migration risk, implementation effort, operational sustainability, and long-term maintainability. Its rationale was to make each architectural change independently verifiable, so that modernization could proceed through controlled increments rather than a single high-risk replacement event.

4.5 Risk Management and Coexistence Plan

Risk management was integrated directly into the migration strategy rather than treated as a separate operational activity.

Migration steps were designed to remain independently verifiable through validation workflows, data-consistency checks, and coexistence-oriented testing mechanisms. Development and production environments were managed separately in order to reduce operational risk during experimentation and rollout phases.

Special attention was given to dual-backend coexistence during data-layer migration stages. Validation workflows were introduced to detect behavioral divergence early and support controlled switchover decisions.

Data replication and validation artifacts were also designed to provide measurable evidence regarding migration progress and consistency verification.

As a consequence, coexistence was treated not as an informal transition phase, but as a structured engineering stage supported by automation, validation, and reproducible infrastructure management practices.

This approach reflects Morris’s principle of reducing the scope of each change as a fundamental risk management practice: smaller increments limit the blast radius of any failure and make it easier to detect, isolate, and correct problems before they propagate [2].

Chapter 5

Target Architecture and Solution Design (To-Be)

This chapter presents the target architecture resulting from the decision framework in Chapter 4. The baseline as-is system is described in Chapter 3; here the focus is on the migration destination designed for incremental adoption.

5.1 Architecture Overview

The migration of Oversight Cloud from the legacy deployment model to AWS required the definition of a cloud-native architecture able to improve scalability, maintainability, and service integration while preserving the original application workflow.

The selected architecture follows a layered model with four levels:

- client and delivery layer;
- load balancing and routing layer;
- application layer;
- managed services layer.

The resulting left-to-right request flow starts from user interaction through the browser, continues through CloudFront and an Application Load Balancer, reaches frontend and backend services on ECS Fargate, and integrates with managed services including PostgreSQL, Redis, Amazon S3, Cognito, SES, and SNS (Figure 5.1).




Figure 5.1. Final AWS architecture adopted for Outsight Cloud.

5.2 Client and Delivery Layer

5.2.1 Browser

The browser is the user entry point to Oversight Cloud. Through the frontend interface, users perform simulation visualization, organization and site management, document access, and file upload/download operations. All interactions occur over HTTPS.

5.2.2 Amazon CloudFront

Amazon CloudFront is used as the content delivery layer between users and the internal application infrastructure. Its role is to reduce latency, cache static assets, accelerate delivery, improve availability, and reduce backend load. This is particularly relevant for frontend static resources generated by the Vue application.

5.3 Load Balancing Layer

5.3.1 Application Load Balancer

The Application Load Balancer (ALB) is the central routing component. It receives requests forwarded by CloudFront and distributes traffic across services running on ECS Fargate. It provides request routing, health checks, HTTPS termination, and traffic distribution, enabling independent scaling and improved fault tolerance.

5.4 Application Layer

5.4.1 ECS Fargate Runtime

The application layer is deployed on Amazon ECS Fargate with a serverless container execution model. Compared to the previous self-managed deployment model, Fargate removes the need to provision, maintain, patch, and scale virtual machines manually.

This runtime choice follows the evaluation described in Chapter 4: among the container orchestration alternatives considered, ECS Fargate offered the best fit for

an incremental migration from the legacy baseline described in Chapter 3. Kubernetes was not selected at this stage because the associated cluster-management overhead (covering networking, upgrades, monitoring, capacity planning, and security governance) was not justified by the scale and operational requirements of Oversight Cloud during the migration.

5.4.2 Frontend Service (Vue)

The frontend service is implemented in Vue.js and is responsible for rendering the user interface, handling user interactions, communicating with backend APIs, and managing authentication-related flows. Deploying it as a dedicated container enables independent release and scaling behavior.

5.4.3 Backend Service (Node.js API)

The backend service is implemented in Node.js and exposes the core APIs used by Oversight Cloud. It orchestrates business logic related to organizations, sites, simulations, authentication, file operations, notifications, and integrations with AWS managed services.

5.5 Managed Services Layer

5.5.1 PostgreSQL

PostgreSQL replaced Airtable as primary data storage. The migration objective was to move from a semi-structured operational backend toward a relational model with stronger consistency and integrity guarantees. Core entities such as users, organizations, sites, and simulations are persisted in PostgreSQL.

5.5.2 Redis

Redis is used as in-memory cache and session storage. It reduces repeated accesses to persistent storage and improves response time for frequently accessed information and temporary state management.

5.5.3 Amazon S3

Amazon S3 provides object storage for simulation files, uploaded assets, documentation artifacts, and deployment-related resources. It was selected for scalability, durability, and native integration with the AWS ecosystem.

5.5.4 Amazon Cognito

Amazon Cognito is used for authentication and identity management, including user registration, login workflows, identity verification, and token handling. This reduces custom authentication logic in the application layer.

5.5.5 Amazon SES

Amazon SES is used for outgoing email communication such as invitations, account-related messages, notifications, and operational communications.

5.5.6 Amazon SNS

Amazon SNS is adopted for asynchronous event-driven notifications. It supports decoupled communication patterns for alerts, internal notifications, workflow events, and service integration triggers.

5.6 Network and Security Topology

The AWS deployment of Oversight Cloud was designed around a simplified cloud-native networking model focused on incremental migration, operational simplicity, and controlled exposure of public services.

The infrastructure operates inside an AWS Virtual Private Cloud (VPC) and relies on multiple public subnets distributed across different Availability Zones in order to improve service availability and fault tolerance. During the migration phase, the selected architecture prioritized deployment simplicity and rapid operational adoption over complex network segmentation strategies.

Application services running on Amazon ECS Fargate are deployed with public outbound connectivity enabled. This configuration simplifies communication with

externally managed services and reduces infrastructure overhead by avoiding additional networking components during the early migration stages.

Despite the use of public subnets, direct external access to application containers is restricted through dedicated Security Groups. Network exposure responsibilities are separated between the Application Load Balancer (ALB) and ECS services: the ALB accepts inbound HTTP and HTTPS traffic from external clients; ECS services accept inbound traffic only from the ALB; application containers are therefore not directly exposed to the Internet.

The Application Load Balancer acts as the primary public entry point for the platform and is responsible for HTTPS termination, request routing, health checks, and traffic distribution across frontend and backend services.

Transport Layer Security (TLS) is managed through AWS Certificate Manager (ACM), enabling centralized certificate lifecycle management and secure HTTPS communication across deployment environments.

The architecture also separates application traffic from static asset delivery. Simulation files and large downloadable resources are stored in private Amazon S3 buckets and distributed through Amazon CloudFront. The CDN layer operates independently from the application load balancer and is dedicated exclusively to content delivery and asset acceleration.

To reduce direct exposure of storage resources, access to S3 objects is restricted through CloudFront Origin Access Control (OAC), allowing content retrieval only through the CDN distribution layer.

Overall, the network topology reflects a pragmatic migration-oriented design that balances security, operational simplicity, and incremental cloud-native adoption while preserving the possibility of future evolution toward more isolated and fully private network architectures.

5.7 Architectural Benefits

Compared to the legacy architecture presented in Chapter 3, the target AWS-based design introduces several architectural improvements that directly address the migration drivers identified in Chapter 1 and the decision criteria discussed in Chapter 4.

First, the architecture provides a clearer separation of concerns between the presentation layer, application services, data storage, and supporting cloud services. This separation reduces coupling between components and simplifies future evolution of individual subsystems without requiring extensive modifications across the platform.

Second, the adoption of AWS managed services significantly reduces operational responsibility. Functions previously handled through self-managed infrastructure, such as load balancing, TLS certificate management, storage scalability, and database availability, are delegated to managed AWS services. This allows engineering effort to focus on application functionality rather than infrastructure maintenance.

The architecture also improves scalability characteristics. Frontend and backend services are deployed independently through ECS Fargate, enabling each component to scale according to its own workload profile. Similarly, managed storage and database services can evolve independently from application deployments, providing greater flexibility as platform demand increases.

From a maintainability perspective, the introduction of Infrastructure as Code through Terraform establishes a reproducible and version-controlled infrastructure model. Infrastructure definitions become reviewable artifacts that can be tracked, validated, and evolved using the same engineering practices applied to application code. This reduces configuration drift and improves consistency across environments.

The migration from Airtable to PostgreSQL further strengthens architectural ownership of the data layer. Explicit relational schemas, database-level constraints, and native query capabilities reduce application-side complexity while improving consistency guarantees and long-term maintainability.

Finally, the target architecture improves operational resilience. Environment separation, automated provisioning, managed services, and the validation mechanisms described in Chapter 6 provide a stronger foundation for controlled deployments, future migrations, and incremental platform evolution.

Chapter 6

Migration Implementation

This chapter describes how the migration strategy was executed in practice, covering the infrastructure transition, data-layer migration, and validation workflows that together enabled the controlled move from the legacy baseline to the target AWS architecture.

6.1 Incremental Execution Plan

The migration was implemented incrementally, structured around four tracks executed in a deliberate sequence:

- establish cloud infrastructure on AWS;
- formalize reproducible provisioning with Terraform;
- migrate data from Airtable to PostgreSQL;
- validate migrated data, application behavior, and deployment correctness before production switchover.

This sequencing made the migration technically traceable: infrastructure readiness, provisioning reproducibility, data-transfer correctness, and application validation were evaluated as separate workstreams, reducing the risk of cascading failures across interdependent components.

The infrastructure migration was executed as a first step, allowing AWS environments, networking components, load balancing, container execution services,

and deployment workflows to be validated independently before the Airtable-to-PostgreSQL data transition took place.

Figure 6.1 summarizes this execution model by showing how the migration moved from the legacy baseline toward the target AWS cloud-native architecture through separate infrastructure, deployment, data-migration, and validation tracks.




Figure 6.1. Incremental execution flow from the legacy baseline to the target AWS cloud-native architecture

6.2 Infrastructure as Code and Environment Separation

6.2.1 Infrastructure as Code (IaC)

Infrastructure as Code (IaC) is a software engineering approach in which infrastructure is provisioned and managed through code instead of manual configuration procedures [7]. The rationale comes from DevOps principles: infrastructure should be treated with the same rigor applied to application code, including versioning, reviewability, traceability, and repeatability.

Under the IaC paradigm, infrastructure definitions are declarative artifacts stored in version control systems. This replaces fragile practices based on manual console operations, ad-hoc scripts, or outdated runbooks. As highlighted by AWS guidance, declarative infrastructure improves consistency and reliability in environment creation and updates [7].

Morris identifies three mutually reinforcing core practices for IaC: defining everything as code, continuously testing and delivering all work in progress, and building small loosely coupled pieces [2]. Code defined as declarative artifacts is reusable, consistent, and transparent: everyone can see how infrastructure is built, changes are auditable, and environments can be recreated deterministically.

A key IaC characteristic is that engineers define the desired final state, while automation tooling computes and executes the required operations. This model provides:

- **Consistency:** reproducible environments across development, testing, and production.
- **Version control:** auditable change history for infrastructure evolution.
- **Automation:** integration with CI/CD-driven deployment processes.
- **Reproducibility:** deterministic creation and update workflows.
- **Maintainability:** reusable infrastructure definitions that evolve over time.

In the Outsight Cloud migration, IaC was adopted to automate AWS provisioning and reduce manual intervention during environment setup and updates.

6.2.2 Terraform-Driven Provisioning

Terraform is an Infrastructure as Code (IaC) tool developed by HashiCorp that enables provisioning and lifecycle management of infrastructure through declarative configuration files [8]. Instead of manually creating resources through cloud consoles, infrastructure is defined as code in human-readable files (HCL), which can be versioned, reviewed, reused, and integrated into software delivery workflows.

Terraform interacts with platforms through providers, i.e., plugins that map Terraform resources to external APIs. In this migration, the AWS provider was used to manage cloud resources including ECS Fargate services, Application Load Balancers, Route 53 DNS records, ACM certificates, VPC networking components, S3 storage, and observability-related infrastructure.

Listing 6.1. Simplified ECS Fargate service definition

```
resource "aws_ecs_service" "backend" {
  name = "outsight-backend"
  cluster = aws_ecs_cluster.main.id
  task_definition = aws_ecs_task_definition.backend.arn

  desired_count = 1
  launch_type = "FARGATE"

  network_configuration {
    subnets = aws_subnet.public[*].id
    security_groups = [aws_security_group.ecs.id]
    assign_public_ip = true
  }
}
```

Listing 6.1 shows a simplified Terraform resource definition used to deploy a backend container as an ECS Fargate service.

6.2.3 Terraform Workflow: Write, Plan, Apply

The Terraform workflow follows three phases [8]:

- **Write:** define the target infrastructure state declaratively in configuration

files.

- **Plan:** generate an execution plan that compares desired state and current state, showing which resources will be created, updated, or removed.
- **Apply:** execute the planned operations while respecting resource dependencies through an internal dependency graph.

This workflow was especially useful for migration safety, because planned changes could be reviewed before execution and applied incrementally.

6.2.4 State Management and Team Collaboration

Terraform keeps a state representation of real infrastructure (`terraform.tfstate`) to compute incremental changes instead of recreating complete environments [8]. In practice, state management enabled drift detection and made environment updates predictable across migration iterations.

For collaborative operations, shared state and controlled execution are essential. This improves consistency across development stages and reduces configuration drift during parallel workstreams.

6.2.5 Rationale for Adopting Terraform

Terraform was selected because it aligned with the migration objectives:

- **Standardization:** infrastructure definitions became versioned artifacts.
- **Automation:** provisioning and updates were integrated into CI/CD-oriented processes.
- **Consistency:** environments were reproducible across development and production contexts.
- **Scalability:** infrastructure changes could be applied incrementally.
- **Maintainability:** modular and declarative definitions reduced manual intervention.

Within the Outsight Cloud project, this translated into a reproducible AWS infrastructure baseline supporting the broader cloudification strategy and reducing operational risk during migration.

6.2.6 Environment Governance

A dedicated effort was made to separate development and production concerns. This included environment-specific configuration handling and controlled evolution of shared resources to avoid production side effects during migration experimentation.

This separation was particularly important during the migration from the legacy infrastructure because it allowed AWS environments to be validated independently before production workloads were progressively redirected to the new platform.

6.3 Infrastructure Migration Execution

The preceding chapters described the legacy platform and the target AWS design, while this section focuses on the concrete execution of the migration itself. The objective was to replace the previous Docker Compose and Traefik deployment model with a reproducible cloud-native foundation without forcing a simultaneous application and data cutover.

6.3.1 Target Environment Definition and Terraform Provisioning

The first implementation step consisted of translating the target architecture into an Infrastructure as Code baseline using Terraform. Rather than manually configuring AWS resources through the console, the platform architecture was defined as a declarative stack covering the core components required to host Outsight Cloud in AWS.

The initial Terraform implementation provisioned the main building blocks of the target platform, including ECS Fargate services for the frontend and backend containers, an Application Load Balancer for public routing, Route 53 records for

DNS management, ACM certificates for HTTPS, CloudWatch log groups for centralized logging, and S3 together with CloudFront for simulation file storage and delivery. This stack established the first operational AWS environment capable of replacing the legacy hosting model.

6.3.2 Routing, Application Adaptation, and Container Deployment

A key difference between the legacy and target deployment models concerned request routing. In the previous environment, Traefik acted as reverse proxy and HTTPS entry point for containerized services. In the AWS architecture, routing responsibilities were transferred to the Application Load Balancer and defined through listener rules that map incoming traffic to the appropriate target groups.

The target design separated public user-interface traffic from API traffic at the load-balancer layer. Rather than depending on proxy-level path rewriting, the back-end service was adapted so that its externally exposed routes aligned with the ALB routing model, while a dedicated health-check endpoint remained available for service and target-group probes. This alignment reduced implicit routing assumptions, simplified the deployment topology, and made request handling easier to validate during migration.

Container images were published through the existing private registry used in the delivery pipeline. To allow ECS tasks to pull images securely, registry credentials were supplied at runtime through AWS Secrets Manager instead of being embedded in Infrastructure as Code artifacts or Terraform state. Frontend and backend workloads were deployed as separate ECS Fargate services, preserving the container packaging model already used in the legacy platform while adapting orchestration and credential handling to AWS operational practices.

6.3.3 Storage, TLS, and Content Delivery

Simulation file storage was redesigned around a private S3 bucket fronted by CloudFront. This replaced the previous S3-compatible storage integration with an AWS-native delivery model that improved scalability and reliability while keeping the

bucket non-public. Access to stored objects was restricted through CloudFront Origin Access Control, and Terraform also defined IAM permissions for upload workflows and cache invalidation.

Public HTTPS access was enabled through ACM-managed certificates integrated with the ALB and CloudFront distribution layers. Route 53 was used to manage DNS records either through an existing hosted zone or through a delegated sub-domain, depending on the deployment environment. Together, these components replaced the certificate and routing responsibilities previously handled by Traefik in the legacy environment.

6.3.4 Environment Validation and Controlled Service Transition

Once the AWS environment was provisioned, validation activities focused on verifying that the new infrastructure behaved correctly before redirecting production traffic. The validation scope included ECS service startup, ALB routing between frontend and backend paths, HTTPS and DNS configuration, centralized logging, and the S3/CloudFront flow for simulation assets.

Development and production concerns were kept separated throughout this process. In particular, experimental infrastructure changes and database modernization activities were isolated in development-oriented environments so that AWS networking, deployment workflows, and service exposure could be tested without affecting the production system still running on the legacy platform.

The final transition step consisted of progressively redirecting public access from the legacy environment to the AWS platform through DNS updates managed in Route 53. This sequencing allowed the team to validate the new stack independently, compare behavior against the existing deployment, and reduce migration risk by avoiding a single disruptive cutover across infrastructure, application routing, and data storage at the same time.

6.4 Data Migration Tooling and Validation

6.4.1 Relational Foundation

The migration path from Airtable to PostgreSQL was implemented with schema and migration tooling designed to preserve compatibility constraints. The objective was not only data transfer, but controlled evolution toward explicit schema ownership.

6.4.2 Replication and Validation Scripts

Dedicated scripts were developed to:

- replicate data from Airtable into PostgreSQL;
- validate consistency through machine-readable reports;
- verify migration correctness at each checkpoint before production switchover.

These artifacts transformed data migration from a one-time operation into a repeatable and verifiable workflow, providing measurable confidence in the correctness of the transition before the final cutover.

Chapter 7

Evaluation and Discussion

7.1 Benchmark Methodology and Demand Profile

The evaluation presented in this chapter is grounded in a reproducible, script-driven benchmark framework developed specifically for this thesis. All measurements are collected through automated HTTP requests against both the legacy and the new AWS environment, with results stored in structured CSV files and post-processed using Python. This section describes the measurement scope, the statistical methodology, and the configuration of both benchmark targets.

7.1.1 Measurement Scope

The benchmark targets the Outsight Cloud web application at the HTTP layer. For each environment, five timing components are extracted from every individual request using `curl`'s built-in timing variables. These components decompose the total round-trip time into the distinct network and application phases, enabling a precise diagnosis of where latency improvements originate. Table 7.1 lists each metric, the corresponding `curl` variable, and its precise definition.

The latency benchmark focuses on the homepage endpoint (`GET /`), which is publicly accessible and representative of the first interaction a user has with the platform. This endpoint exercises the full network stack (DNS resolution, TCP connection, TLS handshake, and application response), making it the most informative

Table 7.1. HTTP timing metrics extracted from `curl` and their meaning. All values are recorded in seconds by `curl` and converted to milliseconds during post-processing.

Metric	curl variable	Description
DNS Lookup	<code>time_namelookup</code>	Elapsed time from request start until DNS name resolution completes.
TCP Connect	<code>time_connect</code>	Elapsed time from request start until the TCP three-way handshake completes.
TLS Handshake	<code>time_appconnect</code>	Elapsed time from request start until the TLS/SSL handshake and key exchange complete.
TTFB	<code>time_starttransfer</code>	Time to First Byte: from request start until the first response byte is received.
Total Time	<code>time_total</code>	Total elapsed time of the complete HTTP transaction, including response body transfer.

single target for an end-to-end comparison. Authenticated API endpoints and static asset downloads were not included in the present analysis, as their evaluation requires a stable authentication token and a consistent asset fingerprint across both environments; the homepage endpoint provides a representative and reproducible baseline for the end-to-end infrastructure comparison.

7.1.2 Workload and Demand Definition

Each benchmark run consists of $N = 100$ independent HTTP GET requests issued sequentially against a single environment. Requests are spaced to avoid connection reuse across measurements: each `curl` invocation opens a fresh TCP connection and performs a complete TLS handshake, ensuring that no warm-up effects inflate the results. Table 7.2 summarises the full set of configuration parameters.

Table 7.2. Benchmark configuration parameters used for all latency measurements.

Parameter	Value / Description
Sample size per environment	$N = 100$ independent requests
Request timeout	30 s (<code>curl -max-time</code>)
Retry policy	1 automatic retry on transport error
HTTP method	GET
Protocol	HTTPS (TLS 1.2 / 1.3)
Confidence level	95 % (Student t -distribution)
Percentiles reported	P90, P95, P99
Measurement tool	<code>curl</code> built-in timing variables
Output format	CSV (one row per request)
Client location	Fixed single location (Paris, FR)

The choice of $N = 100$ balances statistical precision against measurement time. With 100 samples, the 95 % confidence interval for the mean, computed using the Student t -distribution with $n - 1$ degrees of freedom, has a half-width of approximately 3.4 ms for TTFB under the legacy environment (standard deviation ≈ 34 ms), yielding a relative margin of error below 3 %. Increasing N further would reduce this margin proportionally to $1/\sqrt{N}$, offering diminishing returns beyond $N \approx 200$.

In addition to the mean, the benchmark reports the 90th, 95th, and 99th percentiles (P90, P95, P99) of each metric, which are the figures most relevant to

service-level objectives (SLOs). Outliers are not removed: a deliberate design choice that preserves the real-world tail behaviour experienced by end users.

7.1.3 As-Is vs To-Be Benchmark Setup

The two benchmark targets expose the same Outsight Cloud application over HTTPS, but differ substantially in their underlying infrastructure. Table 7.3 compares the two environments across the dimensions most relevant to HTTP latency.

Table 7.3. Infrastructure comparison between the As-Is (Legacy) and To-Be (AWS) benchmark targets. Both environments expose the same application over HTTPS.

Aspect	As-Is	To-Be
Hosting platform	DigitalOcean Droplet	AWS ECS Fargate
Container orchestration	None	ECS
CDN / edge caching	None	Amazon CloudFront
DNS resolver	Standard recursive DNS	Amazon Route 53
TLS termination	Nginx + Let’s Encrypt	CloudFront + ACM certificate
Load balancer	Nginx reverse proxy	AWS Application Load Balancer
Deployment region	Amsterdam	Paris
Measured base URL	<code>legacy.cloud.outsight.ai</code>	<code>cloud.outsight.ai</code>

To ensure a fair comparison, all 200 measurements (100 per environment) were issued from the same client machine located in Paris, France, within the same time window, and with no intermediate caching layer on the client side (`curl -no-buffer -no-keepalive`). The only variable between the two test series is the target base URL, which routes requests to the respective infrastructure. Network conditions on the client side were stable throughout the measurement session, as confirmed by the low variance observed in the DNS and TCP timing components of the AWS environment.

7.2 Cost and TCO Analysis

This section quantifies the economic impact of the migration by comparing the total monthly spend of the legacy platform with that of the new AWS architecture. AWS costs are estimated from the current daily run rate extrapolated to a monthly figure,

thereby excluding the transient overhead incurred during the migration period. All figures are expressed in EUR.

7.2.1 Current Architecture Costs (As-Is)

The legacy Oversight Cloud platform relied on three services:

- **DigitalOcean Spaces: €70/month**, used for object storage of 3D point-cloud datasets and static assets.
- **DigitalOcean Droplets: €102/month**, hosting the entire application stack on manually managed virtual machines.
- **Airtable (7 editor licences): €294/month**, used as the primary operational database and data-management tool.

The total legacy monthly spend was therefore **€466/month**.

7.2.2 Target Architecture Costs (To-Be)

The AWS platform distributes the workload across multiple managed services. Figure 7.1 shows the AWS cost breakdown by service, and Table 7.4 maps each cost category to the corresponding legacy and AWS service. The three largest drivers are **ECS Fargate** (€59/month), **Elastic Load Balancing** (€53/month), and **Amazon RDS** (€48/month). Amazon RDS deserves particular attention: it replaces Airtable as the primary data store while providing full SQL capabilities, automated backups, and high-availability failover, at a fraction of the Airtable licence cost.




Figure 7.1. Estimated steady-state AWS cost by service (EUR/month).

Table 7.4. Monthly cost comparison between the Legacy and AWS platforms.

Category	Legacy	EUR/mo	AWS	EUR/mo
Storage	DigitalOcean Spaces	70	Amazon S3	7.5
Compute	DigitalOcean Droplets	102	ECS Fargate	58.9
Database / SaaS	Airtable (7 editors)	294	Amazon RDS	48.2
Load Balancer	–	–	Elastic LB (ALB)	52.7
Networking	–	–	VPC + Route 53	48.7
Security	–	–	GuardDuty + Secrets Mgr	4.9
Governance	–	–	AWS Config	3.9
Total		466		224.8

7.2.3 Total Cost of Ownership (TCO)

Total cost comparison. Compared to the legacy total of €466/month, the AWS platform reduces the monthly spend to **€231/month**, a saving of **€235/month** (–50%). This result is primarily driven by the elimination of the Airtable licence (€294/month), replaced by Amazon RDS at €48/month.

Category-level breakdown. Figure 7.2 presents the comparison by category. The most striking difference is in the **Database/SaaS** category: Airtable at €294/month is replaced by RDS at €48/month, a reduction of **84 %**. Storage also improves significantly: AWS S3 costs €8/month compared to €70/month for DigitalOcean Spaces (**89 % reduction**), thanks to S3’s pay-per-GB pricing model. Three categories that did not exist in the legacy architecture (*Load Balancer*, *Networking/VPC*, and *Security*) now appear as distinct lines, reflecting services the legacy platform either lacked entirely or embedded invisibly in the Droplet cost.




Figure 7.2. Monthly cost by category: Legacy vs. AWS (EUR/month). Categories absent from the legacy platform are shown as zero.

TCO conclusion. The migration delivers a net monthly saving of **€235/month** while simultaneously upgrading the platform to a production-grade, fully managed architecture. The saving is structural: it stems from eliminating a SaaS tool whose per-seat pricing model was ill-suited to an infrastructure role, and replacing it with a managed database service priced on actual resource consumption. At the same time, the operational effort freed by the transition, quantified in Section 7.4, provides additional indirect savings on engineering time.

7.3 Latency and Performance Analysis

To complement the cost analysis, an HTTP-level latency benchmark was conducted to quantify the end-to-end performance improvement introduced by the cloud migration. The benchmark targets the same service exposed under both the legacy DigitalOcean environment and the new AWS architecture, measuring five timing components extracted directly from `curl`’s built-in timing variables: *DNS resolution*, *TCP connection establishment*, *TLS handshake*, *Time to First Byte* (TTFB), and *Total request time*. A sample of $N = 100$ independent requests was collected for each environment, with a 95% confidence interval (CI) computed from Student’s t -distribution. All measurements were performed from the same client location to ensure comparability. Table 7.5 reports the complete descriptive statistics for both environments.

Table 7.5. Descriptive statistics for the Home endpoint benchmark ($N = 100$ requests per environment). All timing values in milliseconds (ms).

Environment	Metric	Mean	Median	Std Dev	Max	P90	P95	P99
Legacy (DigitalOcean)	DNS Lookup	14.92	7.58	15.38	114.92	30.01	36.39	60.70
	TCP Connect	15.19	7.81	15.45	115.19	30.32	37.47	60.99
	TLS Handshake	84.83	81.40	31.20	333.05	104.09	118.50	169.66
	TTFB	130.74	124.85	34.08	394.91	156.69	177.00	211.58
	Total Time	130.83	124.94	34.09	395.02	156.80	177.13	211.71
AWS (ECS Fargate)	DNS Lookup	1.64	1.15	2.20	22.51	2.43	2.49	2.92
	TCP Connect	1.79	1.27	2.23	22.74	2.67	2.78	3.25
	TLS Handshake	44.45	43.30	13.82	146.38	55.53	58.90	71.90
	TTFB	81.09	78.59	19.73	197.31	100.82	105.22	149.57
	Total Time	81.16	78.64	19.75	197.43	100.92	105.29	149.69

7.3.1 As-Is Latency Measurements

Under the legacy architecture, every component of the HTTP timing stack exhibits both a high mean and significant variance. DNS resolution averages **14.92 ms** (median 7.58 ms), but its standard deviation of 15.38 ms and a 99th-percentile value of 60.70 ms reveal substantial jitter: a fraction of requests suffer DNS lookup times an order of magnitude above the typical case. This variability is attributable to the absence of a dedicated DNS resolver close to the client and the lack of anycast routing in the legacy setup.

TCP connection time closely mirrors the DNS distribution (mean 15.19 ms), since TCP establishment is initiated immediately after name resolution; outliers in one phase cascade into the other. The TLS handshake is the dominant latency contributor in the legacy environment, averaging **84.83 ms** (p95 = 118.50 ms, max = 333.05 ms), a consequence of the server being hosted on a single-region DigitalOcean droplet with no TLS session resumption optimisation.

The aggregate TTFB, which best captures the user-perceived responsiveness of the application, reaches a mean of **130.74 ms** (p95 = 177.00 ms, max = 394.91 ms). A 95% CI width of 13.53 ms confirms that variability is not negligible: some users routinely experience response times more than three times the minimum observed value.

7.3.2 To-Be Latency Measurements

The AWS architecture, built on ECS Fargate behind an Application Load Balancer with CloudFront as the global CDN layer, dramatically reduces every timing component.

DNS resolution drops to a mean of **1.64 ms** (p99 = 2.92 ms), nearly two orders of magnitude lower than the legacy average. This improvement stems from Route 53’s anycast infrastructure, which serves DNS responses from the point-of-presence closest to the client, combined with the short TTLs configured on the CloudFront distribution. The standard deviation falls to 2.20 ms, a ninefold reduction, indicating that DNS resolution is stable and predictable.

TCP connection time follows the same pattern, averaging **1.79 ms**. The TLS handshake settles at a mean of **44.45 ms** (p95 = 58.90 ms), benefiting from CloudFront’s TLS 1.3 support and session-ticket resumption, which avoids the full two-round-trip handshake for returning clients.

Most importantly, TTFB falls to **81.09 ms** on average (p95 = 105.22 ms, max = 197.31 ms), with a standard deviation of 19.73 ms. The CI width narrows to 7.83 ms, indicating a more consistent experience across requests.

7.3.3 Comparative Results

Figure 7.3 presents the mean latency for each component under both environments, together with the 95% CI error bars and the percentage improvement achieved by the migration. The improvement ranges from **−89%** on DNS and TCP (where the legacy architecture paid the highest penalty) to **−38%** on TTFB and Total time. Even though the TLS reduction (−48%) and TTFB reduction (−38%) appear more modest in percentage terms, they translate to absolute savings of 40 ms and 50 ms respectively on *every single request*, which is perceptible to end users.




Figure 7.3. Mean latency comparison, Legacy vs. AWS, per HTTP timing component, with 95% CI error bars.

Figure 7.4 complements the mean comparison by showing the full distribution of each metric. The legacy boxes span a much wider interquartile range and are accompanied by numerous high-value outliers, confirming that the legacy environment is susceptible to occasional latency spikes. By contrast, the AWS boxes are compact and nearly outlier-free for DNS and TCP, and show significantly reduced spread for TLS and TTFB. This tighter distribution is operationally relevant: it implies that the worst-case experience for any given user is now closer to the median, rather than an unpredictable multiple of it.




Figure 7.4. Box plots of latency distributions per HTTP timing component (100 samples per environment).

The empirical Cumulative Distribution Functions (ECDFs) in Figure 7.5 provide a comprehensive view of the tail behaviour. For every metric, the AWS curve reaches high cumulative probability values at latencies well below where the legacy curve begins to plateau, confirming that even the worst-percentile AWS requests outperform the median legacy requests for DNS and TCP. For TTFB, the legacy distribution exhibits a long right tail that extends beyond 350 ms, whereas 99% of AWS requests complete within 150 ms. The p90 and p95 reference lines make it easy to read off the service-level implications directly from the plot: the AWS p95 TTFB (105 ms) is lower than the legacy *median* (125 ms), a result that underscores the magnitude of the performance gain.




Figure 7.5. ECDFs of latency per HTTP timing component, with p90/p95/p99 reference lines.

In summary, the migration to AWS reduces mean end-to-end latency by approximately 38% while simultaneously tightening the distribution, eliminating the long

tail observed on the legacy platform. The gains are most pronounced in the network-proximity-sensitive phases (DNS, TCP), where the combination of Route 53, CloudFront, and the AWS backbone replaces an unoptimised single-region setup. The residual TLS and TTFB latencies are now bounded by the application server processing time and the physical distance between the AWS region and the client, both of which represent architectural lower bounds rather than avoidable inefficiencies.

7.4 Operational Cost and Efficiency

Beyond direct infrastructure spending, the architecture of a platform determines the amount of engineering time required to keep it running. This section analyses the recurring operational tasks associated with the legacy and AWS architectures, and estimates their cost in engineering hours.

7.4.1 As-Is Operational Effort

The legacy platform, built on DigitalOcean Droplets running Docker Compose, placed the full burden of infrastructure operations on the engineering team. The recurring tasks fell into four categories.

Operating system and runtime maintenance. Each Droplet runs a Linux distribution that must be kept up to date. Security patches, kernel updates, and package upgrades must be applied manually, tested against the running application stack, and coordinated across multiple machines to avoid downtime. Docker Engine itself requires version management: breaking changes between Docker releases have historically required coordinated engineering effort across both the host OS and the Compose configuration. On average, OS and runtime maintenance consumed approximately **4–6 hours per month** of engineering time, including the verification of applied patches and monitoring for regressions.

Deployment and rollback management. In the legacy setup, deployments were performed by SSH-ing into the Droplet, pulling the new container images, and restarting services with Docker Compose. This procedure was error-prone and lacked atomicity: a failed service start could leave the system in a partially updated

state, requiring manual inspection and rollback. There was no built-in blue/green or rolling deployment mechanism; any deployment carried a brief service interruption. Engineering time for deployment operations (including preparation, execution, and validation) was estimated at **3–4 hours per deployment cycle**.

Scaling and capacity management. Horizontal scaling on DigitalOcean required manually resizing or duplicating Droplets through the web console or CLI. There was no automatic scaling policy: traffic spikes had to be anticipated and handled reactively. Vertical scaling (increasing CPU or RAM on an existing Droplet) required a reboot, resulting in planned downtime. Capacity planning reviews were conducted periodically and accounted for approximately **2–3 hours per month**.

Monitoring, alerting, and incident response. The legacy platform had no centralised observability stack. Log collection was ad hoc, relying on direct access to Docker container logs via SSH. Setting up and maintaining custom monitoring scripts, uptime checks, and alert notifications was the responsibility of the team. Incident diagnosis (identifying which container had failed, reading the relevant logs, and correlating events) was slow and largely manual, typically adding **1–2 hours per incident** in investigation overhead.

The cumulative recurring effort across all four categories amounted to an estimated **10–15 hours per month** of engineering time devoted to undifferentiated infrastructure maintenance, excluding incident response. At a loaded engineering cost of approximately €80–100/hour, this corresponds to a **recurring operational overhead of €800–1 500/month** that adds no direct product value.

7.4.2 To-Be Operational Effort

The AWS architecture delegates the majority of the operational tasks described above to managed services, fundamentally changing the engineering team’s role from infrastructure operator to configuration manager.

No OS or runtime maintenance. ECS Fargate is a serverless container runtime: the team supplies a container image and a task definition, and AWS is responsible for selecting, patching, and managing the underlying compute hosts. There is no Linux

instance to update, no Docker Engine to maintain, and no kernel to patch. Container runtime upgrades are handled transparently by the platform. The engineering effort previously spent on OS and runtime maintenance is **effectively zero** in the To-Be architecture.

Automated deployments with zero downtime. ECS rolling deployments update task instances incrementally, replacing old containers with new ones while keeping the service available. A deployment is triggered by pushing a new container image to ECR and updating the task definition, a process fully automated through the CI/CD pipeline. Health checks performed by the Application Load Balancer ensure that traffic is only routed to healthy task instances. Rollbacks are instantaneous: re-registering a previous task definition revision reverts the deployment without any manual intervention. Deployment-related engineering effort is reduced to **pipeline maintenance**, which is largely a one-time setup cost.

Automatic scaling. ECS Service Auto Scaling adjusts the number of running Fargate tasks in response to CPU utilisation, memory pressure, or custom CloudWatch metrics, without any manual intervention. Capacity planning is replaced by policy configuration: the team defines minimum, desired, and maximum task counts, and the platform handles runtime adjustments automatically. This eliminates the reactive scaling work that characterised the legacy environment.

Centralised observability. Amazon CloudWatch collects logs from all ECS tasks, ALB access logs, and RDS performance insights in a single pane. Structured log queries, metric dashboards, and threshold-based alarms are configured declaratively as part of the infrastructure-as-code definitions. Incident diagnosis time is reduced significantly: instead of SSH-ing into individual VMs, the team can filter correlated log streams by service, time window, or error pattern within seconds. AWS GuardDuty provides continuous threat detection with no manual configuration or signature updates required.

Database administration. Amazon RDS automates routine database administration tasks including automated backups (point-in-time recovery up to 35 days), minor version upgrades, storage auto-scaling, and failover to a standby replica. The

same tasks on a self-managed PostgreSQL instance would require dedicated DBA effort for backup scheduling, replication monitoring, and patch application.

The residual recurring operational effort in the To-Be architecture is estimated at **2–3 hours per month**, covering configuration updates, cost monitoring, IAM policy reviews, and CI/CD pipeline maintenance.

7.4.3 Operational Cost Comparison

Table 7.6 summarises the operational effort comparison between the two architectures.

Table 7.6. Operational effort comparison between the Legacy and AWS architectures. Estimates are in engineering hours per month.

Activity	Legacy (h/mo)	AWS (h/mo)
OS and runtime patching	4–6	0
Deployment and rollback	3–4	0.5
Scaling and capacity mgmt	2–3	0
Monitoring and log management	2–4	0.5
Database administration	2–3	0.5
Incident investigation overhead	1–2/incident	minimal
Total (excl. incidents)	13–20	1.5–2

At a conservative loaded engineering rate of €80/hour and assuming the mid-point of each range, the legacy platform consumed approximately **€1 320/month** in operational engineering effort. The AWS platform reduces this to approximately **€140/month**, a saving of **€1 180/month** in engineering time alone.

Combined with the €235/month saving on direct infrastructure costs (Section 7.2), the migration yields a total estimated monthly saving of **€1 415/month**, corresponding to an annual saving of approximately **€17 000**. This figure does not account for the reduction in incident frequency and resolution time enabled by centralised observability, which would further improve the return on investment over time.

7.4.4 External Validation of the Operational Estimates

The operational effort figures reported in Table 7.6 are grounded in direct team experience with both environments and therefore represent informed engineering estimates rather than precisely measured timesheet data. To assess whether the *magnitude* of the observed reduction is plausible, it is useful to compare it against independent industry evidence.

A Total Economic Impact (TEI) study conducted by Forrester Consulting and commissioned by AWS quantifies the operational benefits of moving from legacy environments to AWS-managed cloud operations [9]. The study aggregates interviews with five enterprise decision-makers into a single composite organization and reports, among other findings, provisioning time savings of approximately **97.5%** through templated and automated environment creation, a **30%** improvement in DevOps productivity enabled by deployment automation, and a **90%** reduction in high-priority (P1/P2) incidents after the adoption of managed services and Infrastructure-as-Code practices.

Although the Forrester study refers to a large enterprise and to a broad AWS Cloud Operations portfolio rather than to the specific ECS Fargate, RDS, and Terraform stack adopted in this work, the mechanisms responsible for the savings are the same ones that drive the reductions measured here: automated and reproducible provisioning, delegation of patching and scaling to managed services, and centralised observability. The order-of-magnitude reduction in operational effort estimated for Outsight Cloud, from 13–20 to 1.5–2 engineering hours per month (roughly a 90% decrease), is therefore consistent with these independently reported figures, which reinforces the credibility of the estimates despite their approximate nature.

For transparency, it should be noted that the Forrester study is vendor-commissioned and based on a composite rather than an audited organization. Its figures are used here only as an external plausibility check on the direction and magnitude of the operational savings, and not as a direct measurement of the Outsight Cloud platform.

Chapter 8

Conclusions and Future Work

This thesis investigated the modernization of Oversight Cloud through the migration of a production platform from a legacy infrastructure toward a cloud-native architecture on AWS.

The original system relied on self-managed infrastructure and Docker Compose, while depending on several external services, most notably Airtable as the primary operational datastore. Although this architecture enabled rapid development and operational flexibility during the early stages of the product lifecycle, it progressively introduced limitations related to scalability, maintainability, operational governance, and dependency management.

Rather than approaching migration as a simple infrastructure relocation exercise, this work treated cloud adoption as a broader architectural transformation problem involving infrastructure modernization, data-layer evolution, deployment automation, and operational redesign. Particular attention was given to the engineering artifacts that made this transformation controllable, including Terraform definitions, data migration tooling, validation scripts, and environment-specific deployment workflows.

8.1 Key Outcomes

The evaluation presented in Chapter 7 provides quantitative evidence across three dimensions: direct infrastructure cost, end-to-end HTTP latency, and recurring operational effort.

From a cost perspective (Section 7.2), the migration delivers a structural saving driven primarily by the replacement of *Airtable* with *Amazon RDS*, eliminating a per-seat SaaS pricing model that was architecturally ill-suited to a backend infrastructure role. The reduction in object storage costs further reinforces this outcome. These results directly confirm the first migration objective.

From an operational standpoint (Section 7.4), the delegation of patching, scaling, deployment, and database administration to managed AWS services fundamentally reduces the engineering team’s undifferentiated operational burden. The magnitude of this reduction, both in hours and in associated cost, confirms the third migration objective: Infrastructure as Code and managed services substantively lower operational overhead.

From a performance perspective, the HTTP-level benchmark shows consistent improvements across all measured timing components, with the most pronounced gains in the DNS and TCP phases, where the legacy architecture lacked anycast routing and a CDN layer. The tighter latency distribution achieved by the new architecture means that worst-case user experience is now closer to the median rather than an unpredictable multiple of it.

Beyond the measured outcomes, two structural contributions address the remaining migration objectives. The adoption of *Terraform* as the infrastructure management layer established a reproducible, version-controlled, and auditable provisioning baseline, directly addressing the fourth objective of re-establishing ownership over infrastructure and data assets. The incremental migration design, supported by dual-backend coexistence, data replication tooling, validation scripts, and environment separation, addresses the fifth objective of minimizing migration risk and service disruption. The target architecture built on independently scalable managed services addresses the sixth objective of establishing a foundation capable of supporting long-term platform evolution.

8.2 Limitations and Open Questions

The evaluation covers infrastructure cost, HTTP-layer latency, and operational effort estimates. Two aspects of the measurement scope should be acknowledged explicitly.

The latency benchmark targets the publicly accessible homepage endpoint, which provides a reproducible and infrastructure-representative baseline. Authenticated API endpoints and database-intensive workflows were not included in this analysis, as their evaluation requires a stable authentication token and a consistent asset fingerprint across both environments.

The operational effort figures are based on engineering estimates rather than systematically recorded timesheet data. While these estimates are grounded in direct team experience with both environments, they represent informed approximations rather than precisely measured values.

8.3 Future Work

Several directions remain open for subsequent phases of the migration and for continued platform evolution.

The first direction involves extending the benchmark coverage to authenticated API endpoints and database query performance, in order to complement the infrastructure-level measurements with application-layer evidence of the data-layer transition.

A second area involves load testing and scalability validation. The current evaluation characterizes steady-state latency under sequential requests; measuring system behavior under concurrent load would confirm whether ECS Fargate auto-scaling and RDS connection pooling perform as expected under realistic peak-demand conditions.

A third direction concerns the continued evolution of identity and access management. While the migration introduced Amazon Cognito as the authentication layer, additional improvements could further standardize authorization models and strengthen integration consistency across platform components.

Another important direction is the extension of observability capabilities. Enhanced CloudWatch dashboards and performance baselines established after full cutover would provide the empirical foundation for data-driven operational decisions and regression detection.

Future work may also investigate more advanced cloud-native capabilities as platform usage grows, including refined auto-scaling policies, additional resilience mechanisms, and cost optimization strategies informed by actual production usage

patterns.

Finally, the migration principles discussed in this thesis, namely the decision framework, IaC-first provisioning, dual-backend validation, and incremental rollout planning, could be tested in other industrial contexts to assess their broader generalizability beyond the Oversight Cloud case study.

Overall, this thesis demonstrates that legacy-to-cloud-native migration is most effective when treated as a structured engineering process supported by explicit architectural decisions, reproducible infrastructure management, continuous validation, and controlled incremental evolution. The Oversight Cloud case study shows how modernization can be grounded in concrete engineering practices: declarative infrastructure, managed runtime services, relational data modeling, automated validation, and risk-aware rollout planning. The quantitative evaluation confirms that the migration delivered measurable improvements across cost, latency, and operational overhead, validating the architectural decisions made throughout the modernization process.

Bibliography

- [1] S. Newman, *Building Microservices: Designing Fine-Grained Systems*, 2nd ed. O'Reilly Media, 2021.
- [2] K. Morris, *Infrastructure as Code: Dynamic Systems for the Cloud Age*, 2nd ed. O'Reilly Media, 2021.
- [3] Amazon Web Services. «What is cloud migration?» Accessed: 2026-05-17. [Online]. Available: <https://aws.amazon.com/what-is/cloud-migration/>
- [4] Amazon Web Services. «What is iam?» Accessed: 2026-05-31. [Online]. Available: <https://docs.aws.amazon.com/IAM/latest/UserGuide/introduction.html>
- [5] Airtable. «Introduction to airtable basics». Accessed: 2026-05-25. [Online]. Available: <https://support.airtable.com/docs/introduction-to-airtable-basics>
- [6] The Kubernetes Authors. «Overview of kubernetes». Accessed: 2026-05-30. [Online]. Available: <https://kubernetes.io/docs/concepts/overview/>
- [7] Amazon Web Services. «Infrastructure as code». Accessed: 2026-05-17. [Online]. Available: <https://docs.aws.amazon.com/whitepapers/latest/introduction-devops-aws/infrastructure-as-code.html>
- [8] HashiCorp. «What is terraform?» Accessed: 2026-05-17. [Online]. Available: <https://developer.hashicorp.com/terraform/intro>
- [9] Forrester Consulting, «The total economic impact of AWS cloud operations: Cost savings and business benefits enabled by operating on AWS», Forrester Research, Total Economic Impact Study commissioned by AWS, May 2022.