

Migrating Legacy Services to a Cloud-Native Architecture: A Case Study on Oversight Cloud

Giovanni Russo
Politecnico di Torino
Turin, Italy
s343424

Abstract—Legacy platforms often evolve through years of incremental development, producing architectures that are hard to scale, maintain, and modernize. This thesis addresses a concrete engineering problem: migrating Oversight Cloud, a production platform for LiDAR-related workflows, from a legacy service-oriented setup to a cloud-native architecture on Amazon Web Services (AWS). The original system ran on DigitalOcean virtual machines orchestrated with Docker Compose and relied on Airtable as its primary operational database, an arrangement that progressively caused scalability, consistency, and operational-governance problems. Rather than a big-bang rewrite, the thesis proposes and evaluates an incremental migration strategy based on Infrastructure as Code (Terraform), AWS managed services (ECS Fargate, RDS PostgreSQL, S3, CloudFront), and progressive coexistence between legacy and target components. The design is driven by an explicit decision framework that weighs technical and economic trade-offs. A reproducible, script-driven benchmark shows that the migration cuts direct infrastructure cost by 50% (−€235/month), reduces mean HTTP latency by about 38%, and lowers recurring operational effort by roughly one order of magnitude, for a combined estimated saving near €17,000/year.

Index Terms—cloud migration, cloud-native, AWS, Infrastructure as Code, Terraform, ECS Fargate, PostgreSQL, legacy modernization

I. INTRODUCTION

Cloud computing has changed how software systems are designed, deployed, and operated. Managed services, elastic infrastructure, and deployment automation are hard to reproduce in self-managed environments, and many organizations are therefore modernizing legacy applications toward cloud-native architectures. Migrating a production platform, however, is rarely straightforward: legacy systems accumulate architectural dependencies and operational habits that complicate scalability and long-term evolution.

This work studies the migration of *Oversight Cloud* (OC), a platform that manages LiDAR-related workflows, simulations, organizations, and sites. Before the migration, OC ran on DigitalOcean Droplets orchestrated through Docker Compose and depended on several external services, most notably Airtable as its operational datastore. This stack enabled fast early development but, as usage and integrations grew, it showed clear architectural strain.

The central engineering objective is to migrate OC to a managed cloud environment while regaining explicit ownership of core data and infrastructure assets. The main design choice of the whole project is *incrementality*: every architectural change

is made independently verifiable, so modernization proceeds through controlled steps rather than a single high-risk cutover. This principle is the recurring theme of the thesis, and it is what makes the migration safe to run against a live production system.

II. PROBLEM AND BACKGROUND

Oversight Cloud combined self-managed virtual machines with a set of external service dependencies. Airtable, in particular, had grown from a simple datastore into a critical dependency embedded in application logic and operational workflows. From this baseline, four concrete problems follow.

Weak data model. Relationships between organizations, sites, simulations, and permissions were expressed through linked records and lookup fields rather than relational constraints. Consistency depended on application-side logic, and the same logical field could appear as a scalar or an array depending on context, increasing backend complexity.

Operational fragility. Deployments were performed by SSH-ing into Droplets and restarting Docker Compose, with no atomic rollback, no auto-scaling, and no centralized observability. Routine maintenance (OS patching, capacity planning, incident diagnosis) fell entirely on the engineering team.

Cost and coupling. Airtable’s per-seat pricing (€294/month for seven editor licences) was ill-suited to an infrastructure role, while vendor-specific concepts reduced architectural flexibility.

Migration risk. The platform served production traffic, so it could not tolerate a long interruption or an irreversible switch.

Related work frames the solution space. Newman argues against big-bang rewrites and in favour of incremental delivery, the strangler-fig pattern, and parallel runs [1]. Morris describes Infrastructure as Code and the expand-and-contract pattern for changing live systems safely [2]. Cloud migration is commonly classified as rehosting, replatforming, or refactoring [3]. The OC migration is best described as a hybrid: replatforming the runtime (Docker Compose → ECS Fargate) combined with selective refactoring of the data layer (Airtable → PostgreSQL) and the introduction of IaC. Existing guidance, however, tends to treat technical architecture and economic sustainability separately, and provides limited evidence on intermediate validation. This thesis fills that gap with a decision-oriented case study backed by quantitative evaluation.




Fig. 1. Target AWS architecture. Requests flow from the browser through CloudFront and the Application Load Balancer to the Vue and Node.js services on ECS Fargate, which integrate with managed services (RDS PostgreSQL, Redis, S3, Cognito, SES, SNS).

III. SOLUTION DESIGN

The migration design was driven by an explicit decision framework rather than by technology preference. Cloud providers (AWS, Azure, GCP, DigitalOcean, on-premises) and container runtimes (Kubernetes vs. ECS Fargate) were compared against technical and economic criteria: managed services fit, IaC and automation support, operational complexity, ecosystem maturity, and continuity with the existing baseline.

Runtime. Kubernetes was evaluated but rejected for this stage: its cluster-management overhead (networking, upgrades, monitoring, capacity planning, security governance) was not justified by the scale of OC [4], echoing Newman’s caution against premature platform complexity [1]. **Amazon ECS Fargate** was selected because it runs standard Docker containers with no worker nodes to manage, easing the transition from Docker Compose.

Provider. AWS was chosen for the maturity of its managed-service ecosystem, its strong Terraform integration, and its alignment with the company’s broader cloudification direction. The acknowledged trade-off is that AWS can become expensive under specific scaling conditions, but the balance of service quality, flexibility, and migration continuity favoured it.

The resulting target architecture (Fig. 1) follows a layered model: a delivery layer (browser, CloudFront), a routing layer (Application Load Balancer), an application layer (Vue frontend and Node.js backend on ECS Fargate), and a managed-services layer (RDS PostgreSQL, Redis, S3, Cognito, SES, SNS). PostgreSQL replaces Airtable as the primary datastore, restoring an explicit relational schema with database-level integrity. Static assets live in a private S3 bucket exposed only through CloudFront Origin Access Control, while ACM manages TLS centrally. Security groups ensure that containers accept traffic only from the ALB.

IV. IMPLEMENTATION

The migration was executed as four independently verifiable tracks: (i) establish the AWS environment, (ii) formalize provisioning with Terraform, (iii) migrate data from Airtable to PostgreSQL, and (iv) validate data, application behaviour, and deployment before production switchover. Sequencing the work this way kept the migration traceable and limited the blast radius of any single change.

Infrastructure as Code. The whole target stack (ECS services, ALB, Route 53 records, ACM certificates, CloudWatch log groups, and S3/CloudFront) was defined declaratively in Terraform, following Morris’s core practices: define everything as code, test continuously, and build small loosely coupled pieces [2], [5]. The write-plan-apply workflow allowed each change to be reviewed before execution, and shared state enabled drift detection. Registry credentials were injected at runtime through AWS Secrets Manager instead of being embedded in Terraform state.

Coexistence and validation. Routing responsibilities moved from Traefik to the ALB, and the backend was adapted so its external routes matched the ALB model. Dedicated scripts replicated data from Airtable into PostgreSQL and validated consistency through machine-readable reports, turning data migration into a repeatable, verifiable workflow. Development and production concerns were kept separate, and the final switchover was performed gradually through Route 53 DNS updates, avoiding a single disruptive cutover across infrastructure, routing, and storage.

V. EVALUATION

The migration was evaluated with a reproducible, script-driven benchmark along three dimensions: direct infrastructure cost, end-to-end HTTP latency, and recurring operational effort. Latency was measured with `curl` timing variables over $N = 100$ independent requests per environment, issued from the same client in Paris with fresh connections (no keep-alive), and summarized with 95% confidence intervals.

Cost. The legacy platform cost €466/month (DigitalOcean Droplets €102, Spaces €70, Airtable €294). The AWS platform costs €231/month, a **50% reduction** (Fig. 2). The saving is structural: Airtable (€294) is replaced by Amazon RDS (€48), an 84% cut in the database category, while object storage drops 89% (S3 €8 vs. Spaces €70).

Latency. The AWS architecture reduces every HTTP timing component (Fig. 3). DNS and TCP fall by about 89%, thanks to Route 53 anycast and CloudFront points of presence; the TLS handshake drops 48% (85 ms → 44 ms) via TLS 1.3 and session resumption. The user-perceived Time to First Byte (TTFB) falls 38% (131 ms → 81 ms), and the AWS 95th-percentile TTFB (105 ms) is lower than the legacy *median* (125 ms). Beyond the mean, the distribution becomes much tighter: the long tail of the legacy platform (TTFB up to 395 ms) disappears, so the worst-case user experience is now close to the median.

Operational effort. Delegating patching, scaling, deployment, and database administration to managed services reduces




Fig. 2. Monthly cost by category, legacy vs. AWS. Categories absent from the legacy platform (load balancer, networking, security) appear as zero.




Fig. 3. Mean latency per HTTP timing component, legacy vs. AWS, with 95% CI error bars and percentage improvement.

TABLE I
MAIN EVALUATION RESULTS

Metric	Result
Infrastructure cost	€466 → €231/mo (−50%)
Database (Airtable→RDS)	€294 → €48/mo (−84%)
Object storage (Spaces→S3)	€70 → €8/mo (−89%)
Mean TTFB	131 → 81 ms (−38%)
DNS / TCP latency	≈ −89%
Operational effort	13–20 → 1.5–2 h/mo
Combined saving	≈€1,415/mo (≈€17k/yr)

recurring operational work from an estimated 13–20 h/month to 1.5–2 h/month. At a conservative €80/h loaded rate, this is a further €1,180/month saving on engineering time. Combined with the €235/month infrastructure saving, the migration yields roughly €1,415/month, about €17,000/year. This roughly 90% drop in operational effort is consistent with independent industry evidence: a Forrester Total Economic Impact study commissioned by AWS reports around 97.5% provisioning time savings, a 30% DevOps productivity gain, and a 90% reduction in high-priority incidents after adopting managed cloud operations [6]. Table I collects the main results.

VI. CONCLUSION

This thesis showed that a legacy-to-cloud-native migration is most effective when treated as a structured engineering process rather than a simple relocation. An explicit decision framework selected AWS managed services and ECS Fargate

over more complex or less aligned alternatives; an IaC-first, incremental strategy with dual-backend coexistence and automated validation made each step independently verifiable and reversible. The quantitative evaluation confirms measurable gains across all three dimensions: 50% lower infrastructure cost, 38% lower mean latency with a much tighter distribution, and roughly a tenfold reduction in operational effort.

The work also has clear limits. The latency benchmark targets the public homepage endpoint; authenticated API and database-intensive workflows were not measured. Operational figures are informed estimates rather than recorded timesheets, and steady-state latency was measured under sequential rather than concurrent load. These points define natural future work: extending benchmarks to authenticated and database-heavy paths, load and scalability testing of Fargate auto-scaling and RDS pooling, deeper IAM standardization with Cognito, and richer observability. Beyond OC, the decision framework, IaC-first provisioning, dual-backend validation, and incremental rollout planning are transferable to other legacy modernization scenarios.

REFERENCES

- [1] S. Newman, *Building Microservices: Designing Fine-Grained Systems*, 2nd ed. O'Reilly Media, 2021.
- [2] K. Morris, *Infrastructure as Code: Dynamic Systems for the Cloud Age*, 2nd ed. O'Reilly Media, 2021.
- [3] Amazon Web Services, “What is cloud migration?,” 2026. [Online]. Available: <https://aws.amazon.com/what-is/cloud-migration/>
- [4] The Kubernetes Authors, “Overview of Kubernetes,” 2026. [Online]. Available: <https://kubernetes.io/docs/concepts/overview/>
- [5] HashiCorp, “What is Terraform?,” 2026. [Online]. Available: <https://developer.hashicorp.com/terraform/intro>
- [6] Forrester Consulting, “The Total Economic Impact of AWS Cloud Operations,” Total Economic Impact study commissioned by AWS, May 2022.